

Preprocessing and Analysis

BUSI 722: Data-Driven Finance II

Kerry Back, Rice University



Outline

1. Build dataset of features, returns, and targets as before
2. Add industry features
3. Preprocessing: standardize features relative to other stocks at the same date
4. Train, predict, and form portfolios in loop as before
5. Interpret model
 - Feature importances
 - Shapley values
 - Features of best and worst portfolios
6. Evaluate portfolio returns: mean-variance frontiers
7. Train and save



- Build dataset of features, returns, and targets as before
- Add preprocessing of features
 - Features standardized relative to other stocks at the same date
 - Add interactions of features
- Interpret model
 - Feature importances
 - Shapley values
 - Features of best and worst portfolios
- Evaluate portfolio returns
 -

1. Create dataset as before

```
In [1]: import numpy as np
import pandas as pd
from sqlalchemy import create_engine
from sklearn.ensemble import RandomForestRegressor
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("whitegrid")
```

```
In [2]: server = 'fs.rice.edu'
        database = 'stocks'
        username = 'stocks'
        password = '6LAZH1'
        driver = 'SQL+Server'
        string = f"mssql+pyodbc://{username}:{password}@{server}/{database}"
        try:
            conn = create_engine(string + "?driver='SQL+Server'").connect()
        except:
            try:
                conn = create_engine(string + "?driver='ODBC+Driver+18+for+SQL+Server'").connect()
            except:
                import pymssql
                string = f"mssql+pymssql://{username}:{password}@{server}/{database}"
                conn = create_engine(string).connect()
```



```
In [3]: sep_weekly = pd.read_sql(
        """
        select date, ticker, closeadj, closeunadj, volume, lastupdated from sep_w
        where date >= '2010-01-01'
        order by ticker, date, lastupdated
        """,
        conn,
    )
    sep_weekly = sep_weekly.groupby(["ticker", "date"]).last()
    sep_weekly = sep_weekly.drop(columns=["lastupdated"])

    ret = sep_weekly.groupby("ticker", group_keys=False).closeadj.pct_change()
    ret.name = "ret"

    price = sep_weekly.closeunadj
    price.name = "price"

    volume = sep_weekly.volume
    volume.name = "volume"
```



```
In [4]: ret_annual = sep_weekly.groupby("ticker", group_keys=False).closeadj.pct_change(12)
ret_monthly = sep_weekly.groupby("ticker", group_keys=False).closeadj.pct_change(1)
mom = (1 + ret_annual) / (1 + ret_monthly) - 1
mom.name = "mom"
```



```
In [5]: weekly = pd.read_sql(
        """
        select date, ticker, pb, marketcap, lastupdated from weekly
        where date>='2010-01-01'
        order by ticker, date, lastupdated
        """,
        conn,
    )
    weekly = weekly.groupby(["ticker", "date"]).last()
    weekly = weekly.drop(columns=["lastupdated"])

    pb = weekly.pb
    pb.name = "pb"
    marketcap = weekly.marketcap
    marketcap.name = "marketcap"
```



```
In [6]: sf1 = pd.read_sql(
        """
        select datekey as date, ticker, assets, netinc, equity, lastupdated from
        where datekey>='2010-01-01' and dimension='ARY' and assets>0 and equity>0
        order by ticker, datekey, lastupdated
        """,
        conn,
    )
    sf1 = sf1.groupby(["ticker", "date"]).last()
    sf1 = sf1.drop(columns=["lastupdated"])

    # change dates to Fridays
    from datetime import timedelta
    sf1 = sf1.reset_index()
    sf1.date = sf1.date.map(
        lambda x: x + timedelta(4 - x.weekday())
    )
    sf1 = sf1.set_index(["ticker", "date"])
    sf1 = sf1[~sf1.index.duplicated()]

    assets = sf1.assets
    assets.name = "assets"
    netinc = sf1.netinc
    netinc.name = "netinc"
    equity = sf1.equity
    equity.name = "equity"

    equity = equity.groupby("ticker", group_keys=False).shift()
    roe = netinc / equity
```



```
In [7]: df = pd.concat(
    (
        ret,
        mom,
        volume,
        price,
        pb,
        marketcap,
        roe,
        assetgr
    ),
    axis=1
)
df["ret"] = df.groupby("ticker", group_keys=False).ret.shift(-1)
df["roe"] = df.groupby("ticker", group_keys=False).roe.fffll()
df["assetgr"] = df.groupby("ticker", group_keys=False).assetgr.fffll()
df = df[df.price >= 5]
df = df.dropna()

df = df.reset_index()
df.date = df.date.astype(str)
df = df[df.date >= "2012-01-01"]

df["target1"] = df.groupby("date", group_keys=False).ret.apply(
    lambda x: x - x.median()
)
df["target2"] = df.groupby("date", group_keys=False).ret.apply(
    lambda x: 100*x.rank(pct=True)
)
```



2. Add industry features

- Deviations from industry medians: is a stock's ROE high relative to its industry, etc.
- Database includes "famaindustry" which is a classification into 48 industries (including other=almost nothing)



```
In [8]: industries = pd.read_sql(
        """
        select ticker, famaindustry as industry from tickers
        """,
        conn,
    )
    industries["industry"] = industries.industry.fillna("Almost Nothing")
    df = df.merge(industries, on="ticker", how="left")
```

```
In [9]: for x in features:
        df[f"{x}_industry"] = df.groupby(
            ["date", "industry"],
            group_keys=False
        )[x].apply(
            lambda x: x - x.median()
        )

features += [f"{x}_industry" for x in features]
```

3. Preprocessing: standardize at each date



We are predicting relative performance. It makes sense to use relative features: how does a stock compare to other stocks at the same date? There are multiple options:

- standard scaler (subtract mean and divide by std dev)
- quantile transformer (map to normal or uniform distribution)
- rank with pct=True (quantile transformer to uniform distribution)

Here we will rank.




```
In [10]: for f in features:
          df[f] = df.groupby("date", group_keys=False)[f].apply(
              lambda x: x.rank(pct=True)
          )
```

4. Train, predict and form portfolios as before

- If we set `train_freq` to a large number, the loop will only train once. Use trained model to predict at all subsequent dates. Do this only for demonstration.
- Should validate but will use `max_depth=4` and `max_features=6` in the random forest.

```

In [11]: train_years = 4 # num years of past data to use for training
train_freq = 100 # num years between training
target = "target2"
model = RandomForestRegressor(max_depth=4, max_features=6)

years = range(2012+train_years, 2024, train_freq)
df2 = None
for i, year in enumerate(years):
    print(year)
    start_train = f"{year-train_years}-01-01"
    start_predict = f"{year}-01-01"
    if year == years[-1]:
        stop_predict = "2100-01-01"
    else:
        stop_predict = f"{years[i+1]}-01-01"
    past = df[(df.date >= start_train) & (df.date < start_predict)]
    future = df[(df.date >= start_predict) & (df.date < stop_predict)].copy()
    model.fit(X=past[features], y=past[target])
    future["predict"] = model.predict(X=future[features])
    df2 = pd.concat((df2, future))

df2.head()

```

2016

Out[11]:

	ticker	date	ret	mom	volume	price	pb	marketcap
--	--------	------	-----	-----	--------	-------	----	-----------

208	A	2016-01-01	-0.023424	0.527914	0.862227	41.78	0.712047	0.895527	0.
-----	---	------------	-----------	----------	----------	-------	----------	----------	----

```
In [12]: num_stocks = 50

grouped = df2.groupby("date", group_keys=False).predict
starting_from_best = grouped.rank(ascending=False, method="first")
best = df2[starting_from_best <= num_stocks]
best_rets = best.groupby("date", group_keys=True).ret.mean()
best_rets.index = pd.to_datetime(best_rets.index)

starting_from_worst = grouped.rank(ascending=True, method="first")
worst = df2[starting_from_worst <= num_stocks]
worst_rets = worst.groupby("date", group_keys=True).ret.mean()
worst_rets.index = pd.to_datetime(worst_rets.index)

all_rets = df2.groupby("date", group_keys=True).ret.mean()
all_rets.index = pd.to_datetime(all_rets.index)
```

4. Interpret

Find feature importances for last trained model



```
In [13]: importances = pd.Series(  
        model.feature_importances_,  
        index=features  
    )  
importances = importances.sort_values(ascending=False)  
importances.round(2)
```

```
Out[13]: volume_industry      0.27  
mom                        0.25  
volume                    0.18  
roe_industry              0.06  
mom_industry              0.06  
roe                       0.06  
pb                        0.04  
marketcap_industry        0.03  
marketcap                 0.02  
pb_industry               0.02  
assetgr                   0.01  
assetgr_industry          0.01  
dtype: float64
```



Shapley values

- Shapley values are a way of calculating the contribution each feature makes to predictions.
- Values are calculated for each observation (each stock/date).
- Can use any part of the data, but look here at last prediction date.
- First look at the distribution of predictions, then at the contributions.




```
In [14]: last_date = df2.date.max()  
df3 = df2[df2.date==last_date]  
df3.predict.describe().round(3)
```

```
Out[14]: count    2970.000  
mean         50.017  
std           1.087  
min          42.771  
25%          49.488  
50%          50.466  
75%          50.835  
max          51.386  
Name: predict, dtype: float64
```



```
In [15]: import shap


explainer = shap.Explainer(model)
shap_values = explainer(df3[features])
```

Mean absolute Shapley values

- Shapley values are positive or negative, depending on whether a feature is positively or negatively related to the prediction.
- Here we average the absolute Shapley values across observations to see which features are on average most important (like `feature_importances`).



```
In [16]: shap.plots.bar(shap_values)
```




Look at Shapley values across observations

- Look at Shapley values one feature at a time
- Plot the Shapley value across observations as a function of the feature
- Shaded plot at bottom is histogram of the feature



```
In [17]: feature = "volume"  
shap.plots.scatter(shap_values[:, feature])
```



Extract best, worst, and all stocks in last portfolios



```
In [18]: best_last = best[best.date==last_date].copy()
worst_last = worst[worst.date==last_date].copy()
all_last = df2[df2.date==last_date].copy()


best_last["group"] = "best"
worst_last["group"] = "worst"
all_last["group"] = "all"

last = pd.concat((best_last, worst_last, all_last))
```


Compare features of best, worst, and all portfolios



```
In [19]: feature = "volume"  
sns.boxplot(last, x="group", y=feature)  
plt.show()
```



6. Evaluate



Add SPY returns



```
In [20]: import yfinance as yf

spy = yf.download("SPY", start=2017)["Adj Close"]
spy = pd.DataFrame(spy)
spy["date"] = spy.index.map(
    lambda x: x + timedelta(4 - x.weekday())
)
spy = spy.groupby(["date"])["Adj Close"].last()
spy = spy.pct_change()

rets = pd.concat((spy, best_rets, worst_rets), axis=1).dropna()
rets.columns = ["spy", "best", "worst"]
```

```
[*****100%*****] 1 of 1 completed
```



Return statistics



```
In [21]: means = 52 * rets.mean()
stdevs = np.sqrt(52) * rets.std()
rf = 0.05
sharpes = (means - rf) / stdevs
stats = pd.concat((means, stdevs, sharpes), axis=1)
stats.columns = ["mean", "std", "sharpe"]
stats.round(2)
```

Out[21]:

	mean	std	sharpe
spy	0.14	0.18	0.52
best	0.35	0.27	1.14
worst	-0.29	0.31	-1.08

In [22]: `rets.corr().round(2)`

Out[22]:


	spy	best	worst
spy	1.00	0.43	0.36
best	0.43	1.00	0.61
worst	0.36	0.61	1.00

Plot performance



In [23]: `logy = False`

```
(1+rets.best).cumprod().plot(label="Best", logy=logy)
(1+rets.spy).cumprod().plot(label="SPY", logy=logy)
(1+rets.worst).cumprod().plot(label="Worst", logy=logy)
plt.legend()
plt.show()
```



Find frontier of SPY, best, and worst



```
In [24]: from cvxopt import matrix
from cvxopt.solvers import qp

cov = rets.cov()
means = rets.mean()


P = cov
A = np.array(
    [
        means,
        [1., 1., 1.]
    ]
)
P = matrix(P.to_numpy())
q = matrix(np.zeros((3, 1)))
A = matrix(A)

mns = []
vars = []
ports = []
for targ in np.linspace(0, 0.5/52, 50):
    b = matrix(
        np.array([targ, 1]).reshape(2, 1)
    )
    sol = qp(
        P=P,
        q=q,
        A=A,
        b=b
```



```
In [25]: mns = 52 * np.array(mns)
sds = np.sqrt(52*np.array(vars))

plt.plot(sds, mns, label="frontier")
plt.scatter(x=[np.sqrt(52)*rets.spy.std()], y=[52*rets.spy.mean()], label="SPY")
plt.xlabel("Standard Deviation")
plt.ylabel("Expected Return")
plt.legend()
plt.show()
```



Find best portfolio with same risk as SPY



```
In [26]: stdev = np.max([s for s in sds if s <= np.sqrt(52)*rets.spy.std()])
indx = np.where(sds==stdev)[0].item()
mean = mns[indx]
port = ports[indx]
print(port.round(2))
print(f"portfolio expected return is {mean:.1%}")
```

```
spy      0.82
best     -0.05
worst     0.23
dtype: float64
portfolio expected return is 3.1%
```

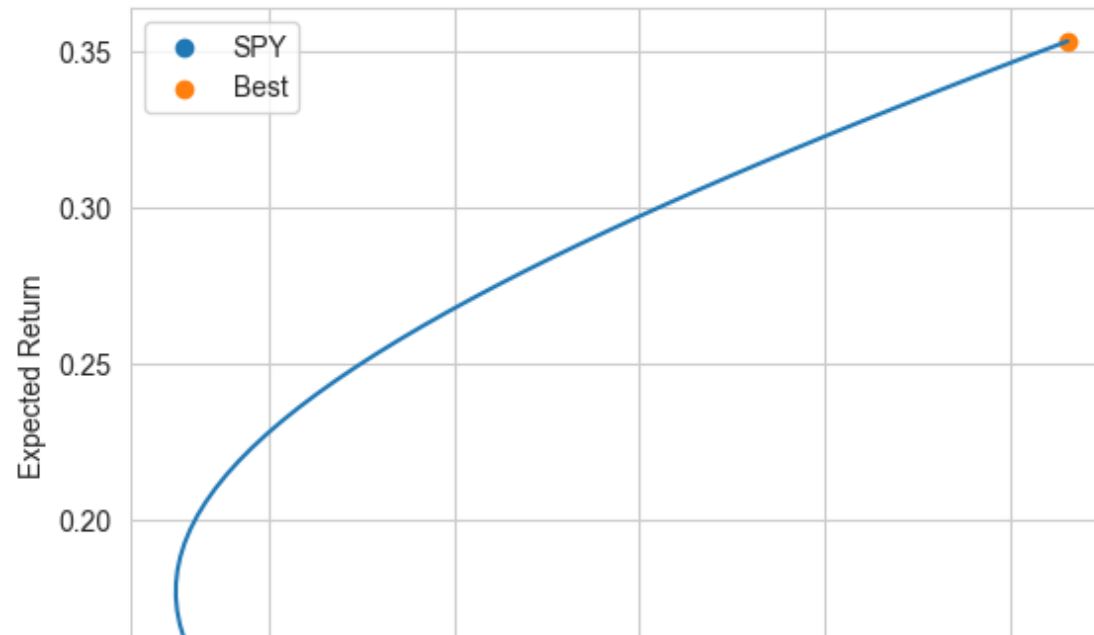


Long-only portfolios of SPY and best




```
In [27]: means = rets[["spy", "best"]].mean()
cov = rets[["spy", "best"]].cov()
ports = [np.array([w, 1-w]) for w in np.linspace(0, 1, 50)]
mns = [52 * means @ w for w in ports]
sds = [np.sqrt(52 * w @ cov @ w) for w in ports]

plt.plot(sds, mns, label=None)
plt.scatter(x=[np.sqrt(52)*rets.spy.std()], y=[52*rets.spy.mean()], label="SPY")
plt.scatter(x=[np.sqrt(52)*rets.best.std()], y=[52*rets.best.mean()], label="Best")
plt.xlabel("Standard Deviation")
plt.ylabel("Expected Return")
plt.legend()
plt.show()
```



140/40 portfolio




```
In [28]: rets["140/40"] = rets.spy + 0.4*rets.best - 0.4*rets.worst
```



```
In [29]: (1+rets[["140/40", "spy"]]).cumprod().plot()
```

```
Out[29]: <AxesSubplot: xlabel='date'>
```



7. Train and save

- Train on the most recent train_years of data
- Save with joblib

```
In [30]: from joblib import dump

         dates = df.date.unique()
         dates.sort()
         date = dates[-52*train_years]
         df3 = df[df.date>=date]
         model.fit(df3[features], df3["target2"])
         dump(model, "mymodel.joblib")
```

```
Out[30]: ['mymodel.joblib']
```

