

Preprocessing and Analysis

BUSI 722: Data-Driven Finance II

Kerry Back, Rice University



Outline

1. Build dataset of features, returns, and targets as before
2. Add preprocessing of features
 - Standardize features relative to other stocks at the same date
 - Add interactions of features
3. Train, predict, and form portfolios in loop as before
4. Interpret model
 - Feature importances
 - Shapley values
 - Features of best and worst portfolios
5. Analyze portfolio returns
 - Mean-variance frontiers of SPY, best, and worst.



- Build dataset of features, returns, and targets as before
- Add preprocessing of features
 - Features standardized relative to other stocks at the same date
 - Add interactions of features
- Interpret model
 - Feature importances
 - Shapley values
 - Features of best and worst portfolios
- Evaluate portfolio returns
 -

1. Create dataset as before

```
In [ ]: import numpy as np
import pandas as pd
from sqlalchemy import create_engine
from sklearn.ensemble import RandomForestRegressor
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("whitegrid")
```

```
In [2]: server = 'fs.rice.edu'
database = 'stocks'
username = 'stocks'
password = '6LAZH1'
driver = 'SQL+Server'
string = f"mssql+pyodbc://{username}:{password}@{server}/{database}"
try:
    conn = create_engine(string + "?driver='SQL+Server'").connect()
except:
    try:
        conn = create_engine(string + "?driver='ODBC+Driver+18+for+SQL+Server'").connect()
    except:
        import pymssql
        string = f"mssql+pymssql://{username}:{password}@{server}/{database}"
        conn = create_engine(string).connect()
```



```
In [3]: sep_weekly = pd.read_sql(
        """
        select date, ticker, closeadj, closeunadj, volume, lastupdated from sep_w
        where date >= '2010-01-01'
        order by ticker, date, lastupdated
        """,
        conn,
    )
    sep_weekly = sep_weekly.groupby(["ticker", "date"]).last()
    sep_weekly = sep_weekly.drop(columns=["lastupdated"])

    ret = sep_weekly.groupby("ticker", group_keys=False).closeadj.pct_change()
    ret.name = "ret"

    price = sep_weekly.closeunadj
    price.name = "price"

    volume = sep_weekly.volume
    volume.name = "volume"
```



```
In [4]: ret_annual = sep_weekly.groupby("ticker", group_keys=False).closeadj.pct_change(12)
ret_monthly = sep_weekly.groupby("ticker", group_keys=False).closeadj.pct_change(1)
mom = (1 + ret_annual) / (1 + ret_monthly) - 1
mom.name = "mom"
```



```
In [5]: weekly = pd.read_sql(
        """
        select date, ticker, pb, marketcap, lastupdated from weekly
        where date>='2010-01-01'
        order by ticker, date, lastupdated
        """,
        conn,
    )
    weekly = weekly.groupby(["ticker", "date"]).last()
    weekly = weekly.drop(columns=["lastupdated"])

    pb = weekly.pb
    pb.name = "pb"
    marketcap = weekly.marketcap
    marketcap.name = "marketcap"
```

```
In [6]: sf1 = pd.read_sql(
        """
        select datekey as date, ticker, assets, netinc, equity, lastupdated from
        where datekey>='2010-01-01' and dimension='ARY' and assets>0 and equity>0
        order by ticker, datekey, lastupdated
        """,
        conn,
    )
    sf1 = sf1.groupby(["ticker", "date"]).last()
    sf1 = sf1.drop(columns=["lastupdated"])

    # change dates to Fridays
    from datetime import timedelta
    sf1 = sf1.reset_index()
    sf1.date = sf1.date.map(
        lambda x: x + timedelta(4 - x.weekday())
    )
    sf1 = sf1.set_index(["ticker", "date"])
    sf1 = sf1[~sf1.index.duplicated()]

    assets = sf1.assets
    assets.name = "assets"
    netinc = sf1.netinc
    netinc.name = "netinc"
    equity = sf1.equity
    equity.name = "equity"

    equity = equity.groupby("ticker", group_keys=False).shift()
    roe = netinc / equity
```



```
In [7]: df = pd.concat(
    (
        ret,
        mom,
        volume,
        price,
        pb,
        marketcap,
        roe,
        assetgr
    ),
    axis=1
)
df["ret"] = df.groupby("ticker", group_keys=False).ret.shift(-1)
df["roe"] = df.groupby("ticker", group_keys=False).roe.fffll()
df["assetgr"] = df.groupby("ticker", group_keys=False).assetgr.fffll()
df = df[df.price >= 5]
df = df.dropna()

df = df.reset_index()
df.date = df.date.astype(str)
df = df[df.date >= "2012-01-01"]

df["target1"] = df.groupby("date", group_keys=False).ret.apply(
    lambda x: x - x.median()
)
df["target2"] = df.groupby("date", group_keys=False).ret.apply(
    lambda x: 100*x.rank(pct=True)
)
```



2. Preprocessing of Features



Features relative to peers

We are predicting relative performance. It makes sense to use relative features: how does a stock compare to other stocks at the same date? There are multiple options:

- standard scaler (subtract mean and divide by std dev)
- quantile transformer (map to normal or uniform distribution)
- rank (quantile transformer to uniform distribution)



```
In [23]: features = [  
    "mom",  
    "volume",  
    "pb",  
    "marketcap",  
    "roe",  
    "assetgr"  
]  
  
for f in features:  
    df[f] = df.groupby("date", group_keys=False)[f].apply(  
        lambda x: x.rank(pct=True)  
    )
```

Interactions

It may be useful to include interactions of features ($x_1 * x_2$ for all features x_1 and x_2). For example, a high value of x_1 may predict a high return only if it is coupled with a high value of x_2 . We could add interactions manually to the dataframe but it is easier to use `PolynomialFeatures`.



Pipeline

We can put preprocessing steps that are to be applied to the entire training set in a pipeline with the model and fit and predict from the pipeline.




```
In [21]: from sklearn.preprocessing import PolynomialFeatures
         from sklearn.pipeline import make_pipeline

         poly = PolynomialFeatures(degree=2, interaction_only=True)
         model = RandomForestRegressor(max_depth=4)
         pipe = make_pipeline(poly, model)
```

3. Train, predict and form portfolios as before

- Only change is to use fit and predict using the pipeline.



```

In [24]: train_years = 5 # num years of past data to use for training
        train_freq = 3 # num years between training
        target = "target2"

        years = range(2012+train_years, 2024, train_freq)
        df2 = None
        for i, year in enumerate(years):
            print(year)
            start_train = f"{year-train_years}-01-01"
            start_predict = f"{year}-01-01"
            if year == years[-1]:
                stop_predict = "2100-01-01"
            else:
                stop_predict = f"{years[i+1]}-01-01"
            past = df[(df.date >= start_train) & (df.date < start_predict)]
            future = df[(df.date >= start_predict) & (df.date < stop_predict)].copy()
            pipe.fit(X=past[features], y=past[target])
            future["predict"] = pipe.predict(X=future[features])
            df2 = pd.concat((df2, future))

        df2.head()

```

2017

```

c:\Users\kerry\AppData\Local\Programs\Python\Python310\lib\site-packa
ges\sklearn\base.py:443: UserWarning: X has feature names, but Random
ForestRegressor was fitted without feature names
  warnings.warn(

```



```
In [10]: num_stocks = 50

grouped = df2.groupby("date", group_keys=False).predict
starting_from_best = grouped.rank(ascending=False, method="first")
best = df2[starting_from_best <= num_stocks]
best_rets = best.groupby("date", group_keys=True).ret.mean()
best_rets.index = pd.to_datetime(best_rets.index)

starting_from_worst = grouped.rank(ascending=True, method="first")
worst = df2[starting_from_worst <= num_stocks]
worst_rets = worst.groupby("date", group_keys=True).ret.mean()
worst_rets.index = pd.to_datetime(worst_rets.index)

all_rets = df2.groupby("date", group_keys=True).ret.mean()
all_rets.index = pd.to_datetime(all_rets.index)
```

4. Interpret

Find feature importances for last trained model



```
In [12]: importances = pd.Series(  
        model.feature_importances_,  
        index=features  
    )  
importances.round(3)
```

```
Out[12]: mom          0.041  
volume      0.304  
pb          0.029  
marketcap   0.170  
roe         0.456  
assetgr     0.000  
dtype: float64
```



Extract best, worst, and all stocks in last portfolios




```
In [13]: last_date = df2.date.max()
best_last = best[best.date==last_date].copy()
worst_last = worst[worst.date==last_date].copy()
all_last = df2[df2.date==last_date].copy()

best_last["group"] = "best"
worst_last["group"] = "worst"
all_last["group"] = "all"


last = pd.concat((best_last, worst_last, all_last))
```

Compare features of best, worst, and all portfolios




```
In [14]: sns.boxplot(last, x="group", y="roe")  
plt.ylim((-4, 4))
```


Out[14]: (-4.0, 4.0)



```
In [15]: sns.boxplot(last, x="group", y="marketcap")  
plt.ylim((0, 0.5e6))  
plt.show()
```



```
sns.boxplot(last, x="group", y="volume")
# plt.ylim((0, 0.5e6))
plt.show()
```



5. Evaluate



Add SPY returns



```
In [17]: import yfinance as yf

spy = yf.download("SPY", start=2017)["Adj Close"]
spy = pd.DataFrame(spy)
spy["date"] = spy.index.map(
    lambda x: x + timedelta(4 - x.weekday())
)
spy = spy.groupby(["date"])["Adj Close"].last()
spy = spy.pct_change()

rets = pd.concat((spy, best_rets, worst_rets), axis=1).dropna()
rets.columns = ["spy", "best", "worst"]
```

```
[*****100%*****] 1 of 1 completed
```




Plot performance



```
In [ ]: logy = False
```

```
(1+rets.best).cumprod().plot(label="Best", logy=logy)  
(1+rets.spy).cumprod().plot(label="SPY", logy=logy)  
(1+rets.worst).cumprod().plot(label="Worst", logy=logy)  
plt.legend()  
plt.show()
```



Find frontier of SPY, best, and worst



```
In [18]: from cvxopt import matrix
from cvxopt.solvers import qp

cov = rets.cov()
means = rets.mean()

P = cov
A = np.array(
    [
        means,
        [1., 1., 1.]
    ]
)
P = matrix(P.to_numpy())
q = matrix(np.zeros((3, 1)))
A = matrix(A)

mns = []
vars = []
ports = []
for targ in np.linspace(0, 0.5/52, 50):
    b = matrix(
        np.array([targ, 1]).reshape(2, 1)
    )
    sol = qp(
        P=P,
        q=q,
        A=A,
        b=b
```



Find best portfolio with same risk as SPY



```
In [19]: stdev = np.max([s for s in sds if s <= np.sqrt(52)*rets.spy.std()])
indx = np.where(sds==stdev)[0].item()
mean = mns[indx]
port = ports[indx]
print(port.round(2))
print(f"portfolio expected return is {mean:.1%}")
```

```
spy      0.53
best      0.92
worst    -0.44
dtype: float64
portfolio expected return is 38.8%
```



Long-only portfolios of SPY and best



```
In [ ]: means = rets[["spy", "best"]].mean()
cov = rets[["spy", "best"]].cov()
ports = [np.array([w, 1-w]) for w in np.linspace(0, 1, 50)]
mns = [52 * means @ w for w in ports]
sds = [np.sqrt(52 * w @ cov @ w) for w in ports]

plt.plot(sds, mns, label=None)
plt.scatter(x=[np.sqrt(52)*rets.spy.std()], y=[52*rets.spy.mean()], label="SPY")
plt.scatter(x=[np.sqrt(52)*rets.best.std()], y=[52*rets.best.mean()], label="Best")
plt.xlabel("Standard Deviation")
plt.ylabel("Expected Return")
plt.legend()
plt.show()
```

